

Abhinita Sanabada

San Jose, CA | (669) 241-8492 | abhinita.sanabada@sjsu.edu | linkedin.com/in/abhinita-sanabada | github.com/Sanabada

Summary

M.S. Data Science candidate leveraging 9 years of software engineering experience in building robust, data-centric systems. Expertise in architecting and automating scalable data pipelines using Airflow, Spark, and Kafka and managing cloud infrastructure and high-availability cloud platforms (AWS/Azure, Docker, Kubernetes) and data-intensive applications.. Seeking a 2026 Data Engineering or MLOps or software engineering internship.

Education

San Jose State University

M.S., Data Science (2025 – Jan 2027, expected)

Coursework: Machine Learning, AI, Deep Learning, Big Data Analytics, Data Visualization, Cloud

Computing, Database Systems, Distributed Systems, Software Engineering, Algorithms

O.E.C

B.E., Electronics and Communication Engineering (2008 - 2012)

Technical Skills

Languages: Python, SQL, Java

Data / ML: Pandas, NumPy, scikit-learn, Matplotlib, Seaborn, statistical analysis

Back-End: Spring Boot, Spring MVC, Microservices, REST APIs

Data Engineering: Airflow, Spark (PySpark), ETL / ETL, data modeling, testing / validation

Platforms: Hadoop, Kafka, Kubernetes, Docker

Cloud / DevOps: AWS, Azure, Git, CI/CD (Jenkins, GitHub Actions), Splunk, Redis

Databases: MySQL, Oracle, MongoDB

Visualization: Tableau, Power BI, Preset

Academic Projects

Real-Time SMS Spam Classification Pipeline

- Architected a real-time, event-driven pipeline to classify SMS messages as 'spam' or 'ham' using a machine learning model.
- The entire system, including Kafka, Zookeeper, and Python-based producers/consumers, was orchestrated using Docker and Docker Compose.
- This project demonstrates a complete, containerized data flow from live ingestion to real-time ML-driven classification.

Scalable Weather Data Ingestion Pipeline

- Built an ML pipeline to predict Architected a scalable data pipeline on HDFS to process raw NOAA-style station feeds, using Apache Spark for batch and streaming ETL. Engineered a real-time ingestion layer using Kafka to handle live data topics. Built a lightweight REST API to serve processed daily/hourly data to downstream applications and a simple GUI (PyQt5) for data validation.
- Models were trained with Spark MLlib (baseline linear/GBTs) and evaluated via time-series splits to avoid leakage.

- I exposed daily/hourly forecasts as Spark jobs writing back to HDFS + a lightweight REST layer, and tracked accuracy with MAE/RMSE for continuous improvements.

Mental Health Risk Prediction (Social & Behavioral Features)

- Designed an AI/ML pipeline to automate the prediction of mental-health risk from social/behavioral signals.
- Trained and compared Logistic Regression and Random Forest with stratified cross-validation, class-imbalance handling, and threshold tuning; evaluated using ROC-AUC/PR-AUC and F1.
- Delivered dashboards (Tableau/Power BI) to explain drivers and model lift, documented assumptions and limits

Professional Experience

Wells Fargo	Senior Software Engineer	Nov 2021–Jan 2024
• Maintained and scaled core banking microservices by building robust REST APIs using Java/Spring, ensuring high availability and transactional integrity. • Created a Selenium + TestNG automation framework from scratch, reducing regression cycles by ~40% • Led monthly releases and triaged post-production issues, driving faster defect resolution and stakeholder updates • Migrated core banking services to Azure and containerized with Docker to improve scalability and performance • <i>Tech:</i> Java (Spring Boot/MVC), SQL, Docker, Azure, Jenkins, JUnit, Unix		
SAP Labs	Software Development Engineer 2	Apr 2019–Oct 2021
• Delivered features for an automotive product on SAP HANA Cloud used by global clients • Built a Kyma/Kubernetes POC that simplified runtime ops and improved traceability across data flows • Managed post-production support and hotfixes using ServiceNow and Splunk to stabilize releases • <i>Tech:</i> Spring Boot, HANA Cloud, GraphQL, Kyma/Kubernetes, Splunk, REST APIs		
J.P. Morgan Chase & Co.	Software Engineer	Jan 2017–Apr 2019
• Built microservices and ingestion pipelines feeding analytics dashboards; improved availability via replication and indexing on Oracle/MongoDB • Implemented data quality checks and metadata practices to support reliable reporting for internal stakeholders • Engineered and automated critical ETL data pipelines using Airflow and Spark, supporting backend analytics systems and stakeholder reporting. • <i>Tech:</i> Python, SQL, Oracle, MongoDB, Airflow, Spark, Tableau, Power BI, CI/CD		
Capgemini	Software Engineer	Apr 2014–Aug 2016
• Built a multi-database abstraction layer (Core Java) to enhance complex transactional systems (Rapid Cash transfers), improving performance and portability across heterogeneous databases. • Delivered JasperReports reporting using Core Java and Oracle SQL		