

Abhinita Sanabada

Sunnyvale, CA | (669) 241-8492 | abhinita.sanabada@gmail.com | linkedin.com/in/abhinita-sanabada | github.com/Sanabada

Summary

Graduate student seeking internships in Data Science, Data Analyst and Data Engineering.

Technical Skills

Languages: Java, Python, SQL

ML: PyTorch, scikit-learn

Databases: PostgreSQL, MySQL, Oracle, MongoDB

Cloud & DevOps: AWS, Docker, Kubernetes, CI/CD, Azure

Tools: Git, Linux, JIRA, Agile/Scrum

Data & Streaming: Apache Kafka, Apache Spark (PySpark), Airflow

Education

San José State University

San Jose, CA

M.S. Data Science

Jan 2025 – May 2027 (expected)

Relevant courses: Machine Learning, Deep Learning, Gen AI, Big Data Analytics, Data Visualization, Cloud Computing, Database Systems, Distributed Systems, Software Engineering, Algorithms. **Eligible** for Summer 2026 internship; **returning to coursework** through May 2027.

Biju Patnaik University of Technology

Btech Electronics and Communication Engineering

Academic Projects

Real-Time SMS Spam Classification Pipeline

- Built a real-time message classification pipeline using Apache Kafka and REST APIs to score incoming texts and support downstream actions.
- Designed a repeatable evaluation and experimentation workflow for model and threshold selection, tracking precision, recall, F1, latency, and drift to keep performance stable over time.
- Implemented data and model versioning and logging to ensure reproducible training runs and auditability of results

Healthcare Affordability Finder System

- **Built a big-data ETL pipeline** on **HDFS + Hadoop MapReduce + PySpark** to ingest CMS hospital price-transparency CSVs (millions of rows), **clean/standardize key fields** (procedure, hospital, gross price, discounted cash) into queryable tables.
- **Trained and evaluated price-prediction models** in **Spark ML** (regression pipelines) to estimate **discounted cash price** from structured features (e.g., gross price, hospital, payer/plan encodings), tracking standard metrics (**RMSE/MAE/R²**).
- **Implemented a semantic search experience** using **SentenceTransformers embeddings** to match user procedure queries (e.g., “knee replacement”, “MRI brain”) and return **cross-hospital price comparisons** via a lightweight UI.

NYC Taxi Price Prediction

- **Built an end-to-end pipeline** for NYC TLC Yellow Taxi pricing analytics: orchestrated **batch taxi ETL** and **near real-time weather ingestion** using **Apache Airflow DAGs**.
- **Implemented scalable PySpark transformations** (schema selection, validation/cleansing) and loaded curated datasets into **Snowflake** via **internal staging + COPY + transactional MERGE** for **idempotent incremental runs**.
- **Delivered an analytics-ready RAW layer** (taxi trips + weather observations) for **dbt/BI** and **prediction feature engineering** by time/zone, with **operational checks** and **robust failure handling**.

Other Experiences & Achievements

Data science intern candidate with hands-on experience building end-to-end data pipelines and machine learning workflows in **Python** and **SQL**, including scalable ETL with **Airflow**, **PySpark**, and **Snowflake**, and model development and evaluation for classification and regression with clear metric tracking; additional engineering depth in **Java/Spring** services and cloud-native tooling (**Docker**, **Kubernetes**, **Azure**), **AWS Fundamentals** certified, and active in the AI

and data community through the **AGI House** building-model competition and **Snowflake** and **NVIDIA** agentic AI conferences.