

Abhinita Sanabada

Sunnyvale, CA | (669) 241-8492 | abhinita.sanabada@gmail.com | linkedin.com/in/abhinita-sanabada | github.com/Sanabada

Summary

Seeking Summer 2026 ML Engineer / Data Analyst / Data Engineer internship,

Technical Skills

Languages: Java, Python, SQL

Data science/ML: Pandas, pytorch, NumPy, scikit-learn, statistics (hypothesis testing, regression, A/B testing), Matplotlib, Seaborn

Data Engineering: Apache Kafka, Apache Spark (PySpark), Airflow, ETL, data modeling

Platforms: Hadoop, HDFS, Apache Kafka, Apache Spark **Databases:** MySQL, Oracle, MongoDB

Cloud/DevOps: AWS, Azure, Snowflake, Docker, Kubernetes, Jenkins, GitHub Actions, CI/CD, Splunk, Redis

Visualization: Tableau, Power BI, Preset **Other:** Git, Linux, JIRA, Agile/Scrum

Education

San José State University

San Jose, CA

M.S., Data Science

Jan 2025 – May 2027 (expected)

Relevant coursework: Machine Learning, Deep Learning, Gen AI, Big Data Analytics, Data Visualization, Cloud Computing, Database Systems, Distributed Systems, Software Engineering, Algorithms

Biju Patnaik University of Technology

Btech Electronics and Communication Engineering

Academic Projects

Real-Time SMS Spam Classification Pipeline

- Built a real-time SMS spam detection service that blocks unwanted and fraudulent texts before they reach users, improving message quality and trust with per-message, low-latency decisions.
- Owned **data/model versioning**, Kept the system accurate and stable via **A/B-style** model/threshold experiments, data/model versioning, and observability on precision/recall/F1, latency, and drift.

Healthcare Affordability Finder System

- Built an end-to-end big data pipeline on HDFS, MapReduce and Spark to ingest **4GB of CMS hospital price data** (13.6M rows), clean and standardize messy multi-hospital files, and produce 47k reliable procedure-level rows for analysis.
- Performed large-scale **EDA and feature selection in PySpark**, showing **heavily right-skewed prices** and a near-fixed discount relative to gross price, then designed a distributed regression pipeline to predict discounted cash prices from gross and payer features.
- Implemented an **NLP semantic search layer** by embedding 47k cleaned procedure descriptions with a **MiniLM SentenceTransformer**, enabling free-text queries (e.g., “knee replacement”) that return relevant procedures along with predicted cash prices.

Stock Prediction using Snowflake

- Built an automated daily **Apache Airflow** pipeline that pulls market **OHLCV** from **yfinance**, standardizes it, and loads via idempotent **MERGE** upserts into **Snowflake**.
- Trained a time-series model with **Snowflake ML** to publish forecasts and actuals to BI-ready tables, owning security/ops (**Airflow Connections/Variables**, raw → model → analytics lineage, safe backfills, reproducibility, auditing).

Other Experience & Achievements

Actively involved in the AI/data community (AGI House building-model competition, Snowflake and NVIDIA agentic AI conferences) and certified in AWS Fundamentals. Please check the link for more details about my hands on experiences.