

Abhinita Sanabada

Sunnyvale, CA | (669) 241-8492 | abhinita.sanabada@gmail.com | linkedin.com/in/abhinita-sanabada | github.com/Sanabada

Summary

Graduate student seeking internships in software engineering and applied AI/ML.

Technical Skills

Languages: Java, Python, SQL

ML: PyTorch, scikit-learn

Backend & APIs: Data Structure & Algorithm, Spring Boot, Spring MVC, Microservices, REST APIs

Databases: PostgreSQL, MySQL, Oracle, MongoDB

Cloud & DevOps: AWS, Docker, Kubernetes, CI/CD, Azure

Testing & Quality: JUnit, Selenium, TestNG

Tools: Git, Linux, JIRA, Agile/Scrum

Data & Streaming: Apache Kafka, Apache Spark (PySpark), Airflow

Education

San José State University

San Jose, CA

M.S. Data Science

Jan 2025 – May 2027 (expected)

Relevant courses: Machine Learning, Deep Learning, Gen AI, Big Data Analytics, Data Visualization, Cloud Computing, Database Systems, Distributed Systems, Software Engineering, Algorithms. **Eligible** for Summer 2026 internship; **returning to coursework** through May 2027.

Biju Patnaik University of Technology

Btech Electronics and Communication Engineering

Academic Projects

Real-Time SMS Spam Classification Pipeline

- Built a scalable real-time message processing backend using Apache Kafka and RESTful APIs.
- Owned data/model versioning, Kept the system accurate and stable via A/B-style model/threshold experiments, data/model versioning, and observability on precision/recall/F1, latency, and drift

Healthcare Affordability Finder System

- **Built a big-data ETL pipeline** on **HDFS + Hadoop MapReduce + PySpark** to ingest CMS hospital price-transparency CSVs (millions of rows), **clean/standardize key fields** (procedure, hospital, gross price, discounted cash) into queryable tables.
- **Trained and evaluated price-prediction models** in **Spark ML** (regression pipelines) to estimate **discounted cash price** from structured features (e.g., gross price, hospital, payer/plan encodings), tracking standard metrics (**RMSE/MAE/R²**).
- **Implemented a semantic search experience** using **SentenceTransformers embeddings** to match user procedure queries (e.g., “knee replacement”, “MRI brain”) and return **cross-hospital price comparisons** via a lightweight UI.

NYC Taxi Price Prediction

- **Built an end-to-end pipeline** for NYC TLC Yellow Taxi pricing analytics: orchestrated **batch taxi ETL** and **near real-time weather ingestion** using **Apache Airflow DAGs**.
- **Implemented scalable PySpark transformations** (schema selection, validation/cleansing) and loaded curated datasets into **Snowflake** via **internal staging + COPY + transactional MERGE** for **idempotent incremental runs**.
- **Delivered an analytics-ready RAW layer** (taxi trips + weather observations) for **dbt/BI** and **prediction feature engineering** by time/zone, with **operational checks** and **robust failure handling**.

Engineering Experience & Achievements

Software engineering intern with experience in **Java/Spring** backend services, test automation (**Selenium/TestNG**), and cloud-native development on **Docker, Kubernetes**, and **Azure**; **AWS Fundamentals** certified and actively involved in the AI/data community through the **AGI House** building-model competition and **Snowflake** and **NVIDIA** agentic AI conferences.